

With fewer than 500 North Atlantic right whales left in the world's oceans, knowing the health and status of each whale is integral to the efforts of researchers working to protect the species from extinction.

Currently, only a handful of very experienced researchers can identify individual whales on sight while out on the water. We aim to build a classifier that use aerial images of whales to predict the individual whales.

```
In [4]: 1 import os
2 import glob
3 import pandas as pd
4 import numpy as np
5 import cv2
6 import matplotlib.pyplot as plt
7 import matplotlib.patches as patches
8 from skimage.feature import hog
9 from numpy import linalg
10 import numpy.matlib
11 from IPython.display import clear_output
12 from time import sleep
```

Dataset Preparation

```
In [2]: 1 all_whales_df = pd.read_csv("train.csv")
```

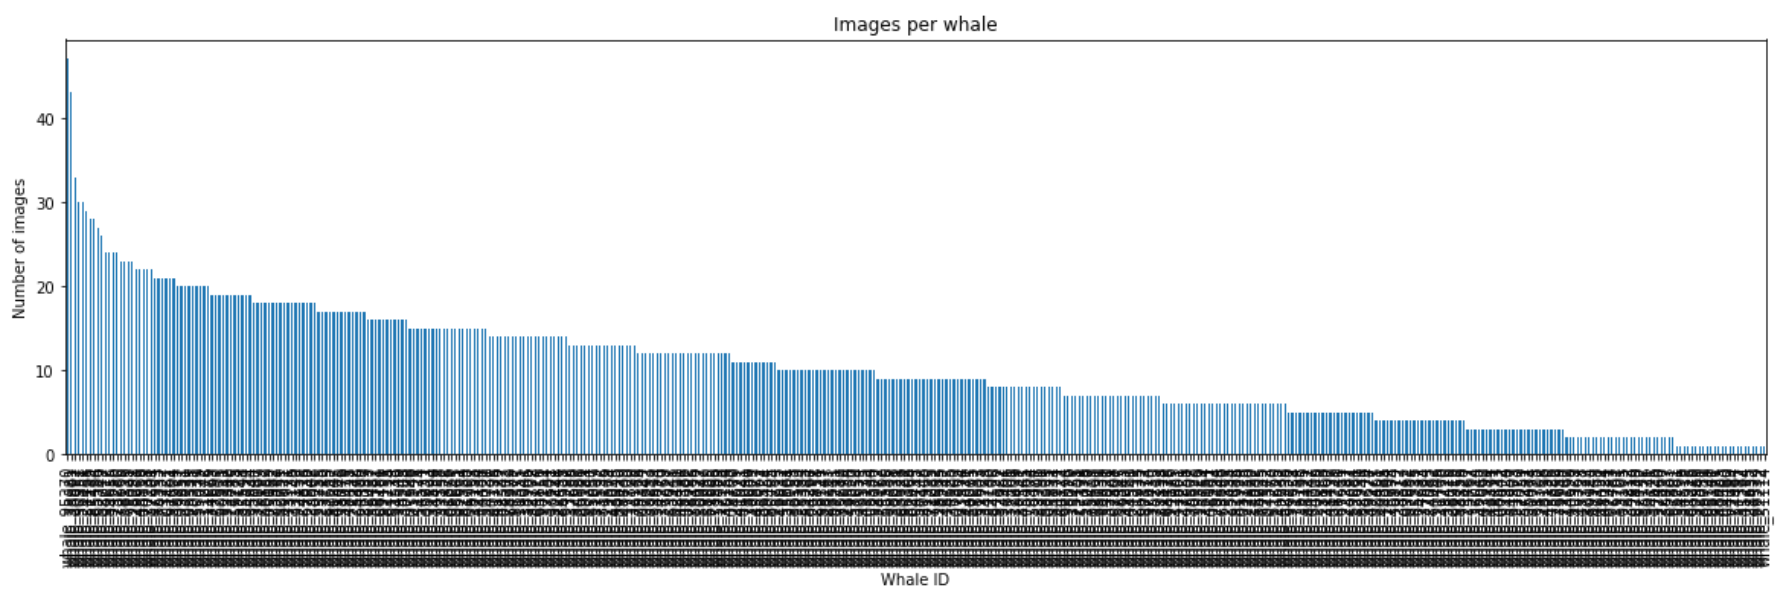
```
In [3]: 1 all_whales_df.head(3)
```

```
Out[3]:
```

	Image	whaleID
0	w_7812.jpg	whale_48813
1	w_4598.jpg	whale_09913
2	w_3828.jpg	whale_45062

```
In [31]: 1 from matplotlib.pyplot import figure
2 import seaborn as sns
3
4 plt.figure(figsize=(20, 5))
5 plt.title('Images per whale')
6 plt.ylabel('Number of images')
7 plt.xlabel('Whale ID')
8
9 # print(list(all_whales_df['whaleID'].value_counts()))
10 all_whales_df['whaleID'].value_counts().plot.bar()
```

```
Out[31]: <AxesSubplot:title={'center':'Images per whale'}, xlabel='Whale ID', ylabel='Number of images'>
```



There are a of whales which less than 10 images in the dataset. So we decide to focus of just five whales where each whale has more than 30 images in the dataset. This gives us enough data for each class of the classifier.

```
In [12]: 1 # sample image from dataset
2 im = cv2.imread('imgs/whale_08017/w_1194.jpg')
3 plt.figure(figsize=(15,15))
4 plt.imshow(cv2.cvtColor(im, cv2.COLOR_BGR2RGB))
5 plt.axis('off')
6 plt.show()
```



The feature that I am most interested in is the callosity pattern which includes the facial markings on top of the head of the whale and the white markings above the blowholes. These features are unique to each whale. According to expert 'Christin B. Khan' on Kaggle(<https://www.kaggle.com/competitions/noaa-right-whale-recognition/discussion/16172> (<https://www.kaggle.com/competitions/noaa-right-whale-recognition/discussion/16172>)), "The white markings on top of the whale's head are called callosities and they change very slightly over time". The dorsal fins, side flippers, and tail shape were other characteristics I chose to disregard even though they would have been informative but would have made detection more difficult.

For this project, we only focus on the marks on the top of the head.


```
In [14]: 1 # The bounding box shows the area we would be focusing on
2 plt.figure(figsize=(15,15))
3 plt.imshow(cv2.cvtColor(im, cv2.COLOR_BGR2RGB))
4 plt.axis('off')
5 rect = patches.Rectangle((1131.3410981697173, 763.3144758735441), 402.1031613976704, 224.905, linewidth=2)
6 plt.gca().add_patch(rect)
7 plt.show()
```



There are annotations available online for the whale heads. We use one of those files with annotations.

```
In [15]: 1 annotations_df = pd.read_json('/Users/pratikaher/FALL22/CV/Whale/right_whale_hunt/annotations/whale_fa
```

```
In [16]: 1 annotations_df.head(2)
```

Out[16]:

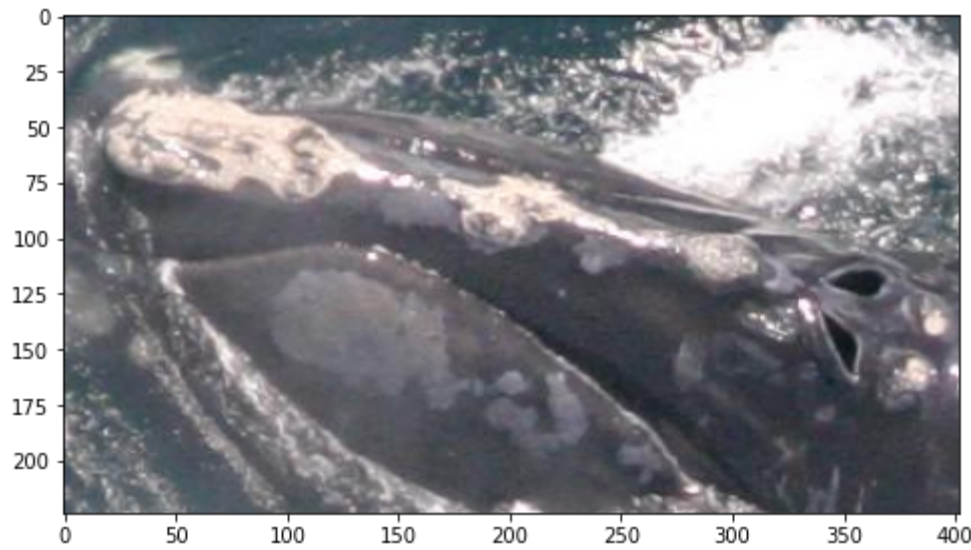
	annotations	class	filename
0	[{'class': 'Head', 'height': 512.0, 'type': 'r...	image	imgs/whale_08017/w_1106.jpg
1	[{'class': 'Head', 'height': 224.9051580698835...	image	imgs/whale_08017/w_1194.jpg

```
In [17]: 1 annotations_df['annotations'][1][0]
```

Out[17]: {'class': 'Head',
'height': 224.90515806988356,
'type': 'rect',
'width': 402.1031613976704,
'x': 1131.3410981697173,
'y': 763.3144758735441}

```
In [18]: 1 def get_cropped_image(im, annotations):
2     y = int(annotations['y'])
3     x = int(annotations['x'])
4     h = int(annotations['height'])
5     w = int(annotations['width'])
6
7     crop_img = cv2.cvtColor(im, cv2.COLOR_BGR2RGB)[y:y+h, x:x+w]
8     return crop_img
```

```
In [19]: 1 plt.figure(figsize=(8,8))
2 crop_img = get_cropped_image(im, annotations_df['annotations'][1][0])
3 plt.imshow(crop_img)
4 plt.show()
```



The above example image is what we will focus on for every whale. For this project, we take the five whales with the most number of images and try to explore various algorithms and techniques to classify those individuals. For those five whales, we will extract the coordinates of the head and crop the images and store them separately. Below is the code that does that.

```
In [66]: 1 path = '/Users/pratikaheer/FALL22/CV/Whale/whale_samples/'
2 write_path = '/Users/pratikaheer/FALL22/CV/Whale/whale_samples/cropped/'
```

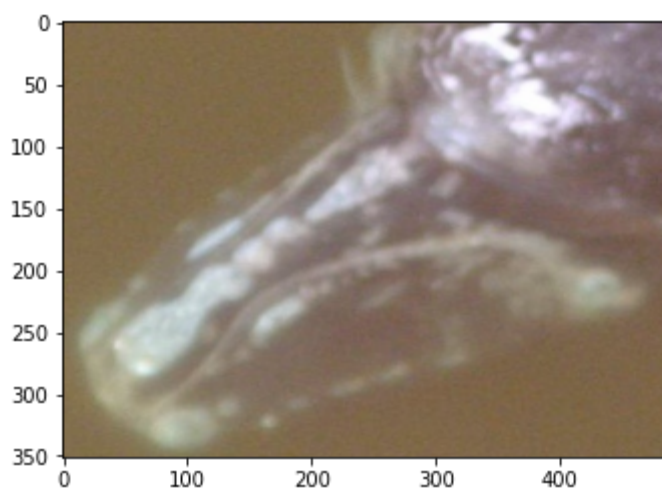
```
In [69]: 1 for file in os.listdir(path):
2     if file != '.DS_Store':
3
4     # print(file, annotations_df[annotations_df['filename'].str.contains(file)])
5     for index, row in annotations_df[annotations_df['filename'].str.contains(file)].iterrows():
6
7         annotations_list = row['annotations']
8
9         if not annotations_list:
10            continue
11
12        for r in annotations_list:
13            if r['class'] == 'Head':
14                annotations = r
15
16        im = cv2.imread(path+file+ '/' +row['filename'].split('/')[-1])
17
18        cropped_image = get_cropped_image(im, annotations)
19
20        # rect = patches.Rectangle((annotations['x'], annotations['y']), annotations['width'], ann
21        # plt.gca().add_patch(rect)
22        # plt.show()
23        # plt.figure(figsize=(15,15))
24        # plt.imshow(cropped_image)
25
26        newpath = write_path + file
27        if not os.path.exists(newpath):
28            os.makedirs(newpath)
29
30        cv2.imwrite(newpath + '/' +row['filename'].split('/')[-1], cv2.cvtColor(cropped_image, cv
31        # plt.imshow(cv2.cvtColor(cv2.cvtColor(cropped_image, cv2.COLOR_BGR2RGB), cv2.COLOR_BGR2RG
32        # plt.show()
```

Orientation seems to be quite important for the classifier to be able to detect images. So we rotate the image so that all images are oriented from top to bottom. We use cv2.findContours and then get the biggest contour and rotate the image around that contour. We assume the biggest contour is the whale head. This helps all images to be aligned from the top-bottom direction. This seems to work in about 80% of images. For the rest, I manually flip the images so that they are aligned.

```
In [70]: 1 countour_example = cv2.imread('/Users/pratikaheer/FALL22/CV/Whale/whale_samples/cropped/whale_48813/w_3
2 countour_example_hsvImage = cv2.cvtColor(countour_example, cv2.COLOR_BGR2HSV)
3 countour_example_gray = cv2.cvtColor(countour_example_hsvImage, cv2.COLOR_BGR2GRAY)
4
5 contours, hierarchy = cv2.findContours(countour_example_gray, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_NONE)
6
7 best = None
8 best_size = 0
9 for cnt in contours:
10    # print(cnt)
11    x,y,w,h = cv2.boundingRect(cnt)
12    size = w * h
13    if size > best_size:
14        best_size = size
15        best = cnt
```



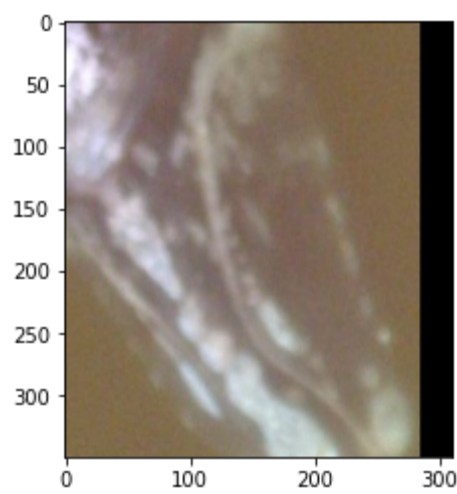
```
In [71]: 1 plt.imshow(countour_example)
          2 plt.show()
```



```
In [72]: 1 ellipse = cv2.fitEllipse(best)
```

```
In [73]: 1 def intround(floatNum):
          2
          3     # print(type(floatNum), isinstance(floatNum, np.ndarray))
          4     if isinstance(floatNum, np.ndarray):
          5         return int(round(floatNum[0]))
          6
          7     return int(round(floatNum))
          8
          9 def rotateAndCrop(img, ellipse):
          10     # Center of rotation is the center of the image
          11     center=tuple(np.array(img.shape[0:2])/2)
          12
          13     # Affine transform - rotation about the center
          14     rotMat = cv2.getRotationMatrix2D(center, ellipse[2], 1.0)
          15
          16     # Recalculate the size of the rotated image
          17     # so the corners do not get cut off
          18     (h, w) = img.shape[0:2]
          19     r = np.deg2rad(ellipse[2])
          20     w, h = (abs(np.sin(r)*h) + abs(np.cos(r)*w), abs(np.sin(r)*w) + abs(np.cos(r)*h))
          21
          22     rotated = cv2.warpAffine(img, rotMat, dsize=(int(w), int(h)), flags=cv2.INTER_LINEAR)
          23
          24     # Get the center of the ellipse and rotate it with the affine transform
          25     centEll = np.ones((3,1))
          26     centEll[0] = ellipse[0][0]
          27     centEll[1] = ellipse[0][1]
          28     centEll = np.dot(rotMat, centEll)
          29
          30     # Get the half lengths of the two axes
          31     mjAxis = max(ellipse[1]) / 2.0
          32     miAxis = min(ellipse[1]) / 2.0
          33
          34     p1x = intround(centEll[0] - miAxis)
          35     p1y = intround(centEll[1] - mjAxis)
          36     p2x = intround(centEll[0] + miAxis)
          37     p2y = intround(centEll[1] + mjAxis)
          38
          39     if p1x < 0:
          40         p1x = 0
          41     if p1y < 0:
          42         p1y = 0
          43     if p2x >= img.shape[1]:
          44         p2x = img.shape[1]-1
          45     if p2y >= img.shape[0]:
          46         p2y = img.shape[0]-1
          47
          48     return rotated[p1y:p2y, p1x:p2x]
          49
```

```
In [74]: 1 plt.imshow(rotateAndCrop(countour_example, ellipse))
2         plt.show()
```



Divide data into train-test-validation split

```
In [4]: 1 cropped_path = '/Users/pratikaheer/FALL22/CV/Whale/whale_samples/cropped/'
2
3 image_names = []
4 target_values = []
5 target_index = 0
6
7 for more_dirs in os.listdir(cropped_path):
8
9     if more_dirs == '.DS_Store':
10         continue
11
12     list_images_in_folder = list(glob.glob(cropped_path+ more_dirs + '/*.jpg'))
13     image_names.extend(list_images_in_folder)
14     target_values.extend([target_index]*len(list_images_in_folder))
15     target_index += 1
```

```
In [6]: 1 image_data = list(zip(image_names, target_values))
```

```
In [7]: 1 import random
2         random.shuffle(image_data)
3
4         test_size = int(len(image_data) * 0.2)
5         valid_size = int(len(image_data) * 0.1)
6         train_size = len(image_data) - test_size - valid_size
7
8         train_data = image_data[:train_size]
9         valid_data = image_data[train_size:train_size+valid_size]
10        test_data = image_data[-test_size:]
```

Feature Exploration

Histogram of oriented gradients feature

```
In [20]: 1 fd, hog_image = hog(cv2.cvtColor(crop_img, cv2.COLOR_BGR2RGB), orientations=8, pixels_per_cell=(8, 8),
2                        cells_per_block=(1, 1), visualize=True, channel_axis=-1)
3
4         plt.figure(figsize=(6,6))
5         plt.axis("off")
6         plt.imshow(hog_image, cmap="gray")
```

```
Out[20]: <matplotlib.image.AxesImage at 0x7fb12320a860>
```



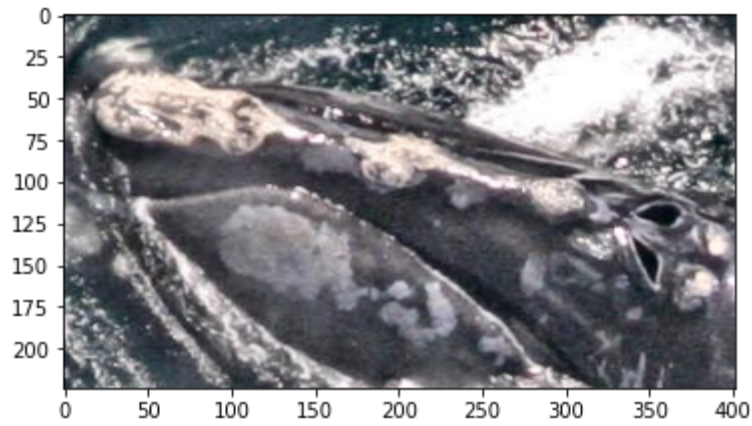
The histogram of gradients was not able to capture the callosity pattern at the top of the head of the whale. This is probably because of a lot of noise at the head region of the whale due to glare of sun and water splashes.

Increase Contrast

I thought increasing the contrast of the image will give us better results as it will make the dark regions darker and bright regions brighter. So we can better see the head.

```
In [21]: 1 clahe = cv2.createCLAHE(clipLimit=2., tileGridSize=(8,8))
          2
          3 lab = cv2.cvtColor(crop_img, cv2.COLOR_BGR2LAB) # convert from BGR to LAB color space
          4 l, a, b = cv2.split(lab) # split on 3 different channels
          5
          6 l2 = clahe.apply(l) # apply CLAHE to the L-channel
          7
          8 lab = cv2.merge((l2,a,b)) # merge channels
          9 contrast_image = cv2.cvtColor(lab, cv2.COLOR_LAB2BGR)
```

```
In [22]: 1 plt.imshow(contrast_image)
          2 plt.show()
```



Visually it did not seem to have that much of a difference.

Changing white pixels to eliminate glare

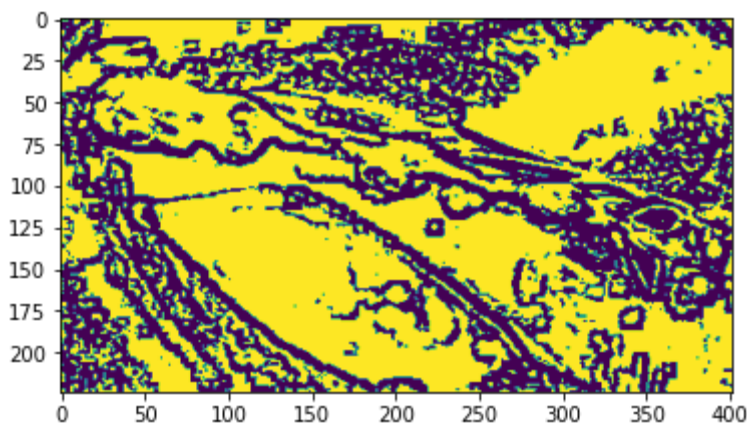
I thought that changing white pixels to black in the image would result in reducing the glare and waves, but it turns out that a lot of white pixels in the trunk are as white or whiter than the glare. So this method would not work as expected.

```
In [257]: 1 gray = cv2.cvtColor(cropped_image, cv2.COLOR_BGR2GRAY)
          2
          3 plt.figure(figsize=(6,6))
          4 ret, thresh = cv2.threshold(gray, 240, 255, cv2.THRESH_BINARY)
          5
          6 cropped_image[thresh == 255] = 0
          7 cropped_image[thresh == 254] = 0
          8 cropped_image[thresh == 253] = 0
          9
          10 plt.imshow(cropped_image)
          11 plt.axis("off")
          12
          13 plt.show()
```



Morphological transformation to remove noise

```
In [268]: 1 image=cv2.cvtColor(crop_img,cv2.COLOR_BGR2GRAY)
2 se=cv2.getStructuringElement(cv2.MORPH_RECT , (8,8))
3 bg=cv2.morphologyEx(image, cv2.MORPH_DILATE, se)
4 out_binary=cv2.threshold(image, 0, 255, cv2.THRESH_OTSU )[1]
5
6 plt.imshow(out_binary)
7
8 plt.show()
```



It needs two inputs, one is our original image, second one is called structuring element or kernel which decides the nature of operation. Two basic morphological operators are Erosion and Dilation. It did not work quite well as there was still noise in the image after morphological transformation

SIFT Features

```
In [275]: 1 import numpy as np
2 import cv2 as cv
3 from matplotlib import pyplot as plt
4
5 MIN_MATCH_COUNT = 4
6
7 img1 = cv.imread('/Users/pratikaheer/FALL22/CV/Whale/investigation/24458/cropped2/w_4611.jpg', 0) # qu
8 img2 = cv.imread('/Users/pratikaheer/FALL22/CV/Whale/investigation/24458/cropped2/w_5271.jpg', 0) # tr
9
10 # Initiate SIFT detector
11 sift = cv.SIFT_create()
12 # find the key points and descriptors with SIFT
13 kp1, des1 = sift.detectAndCompute(img1, None)
14 kp2, des2 = sift.detectAndCompute(img2, None)
```

(181, 342) (478, 226)

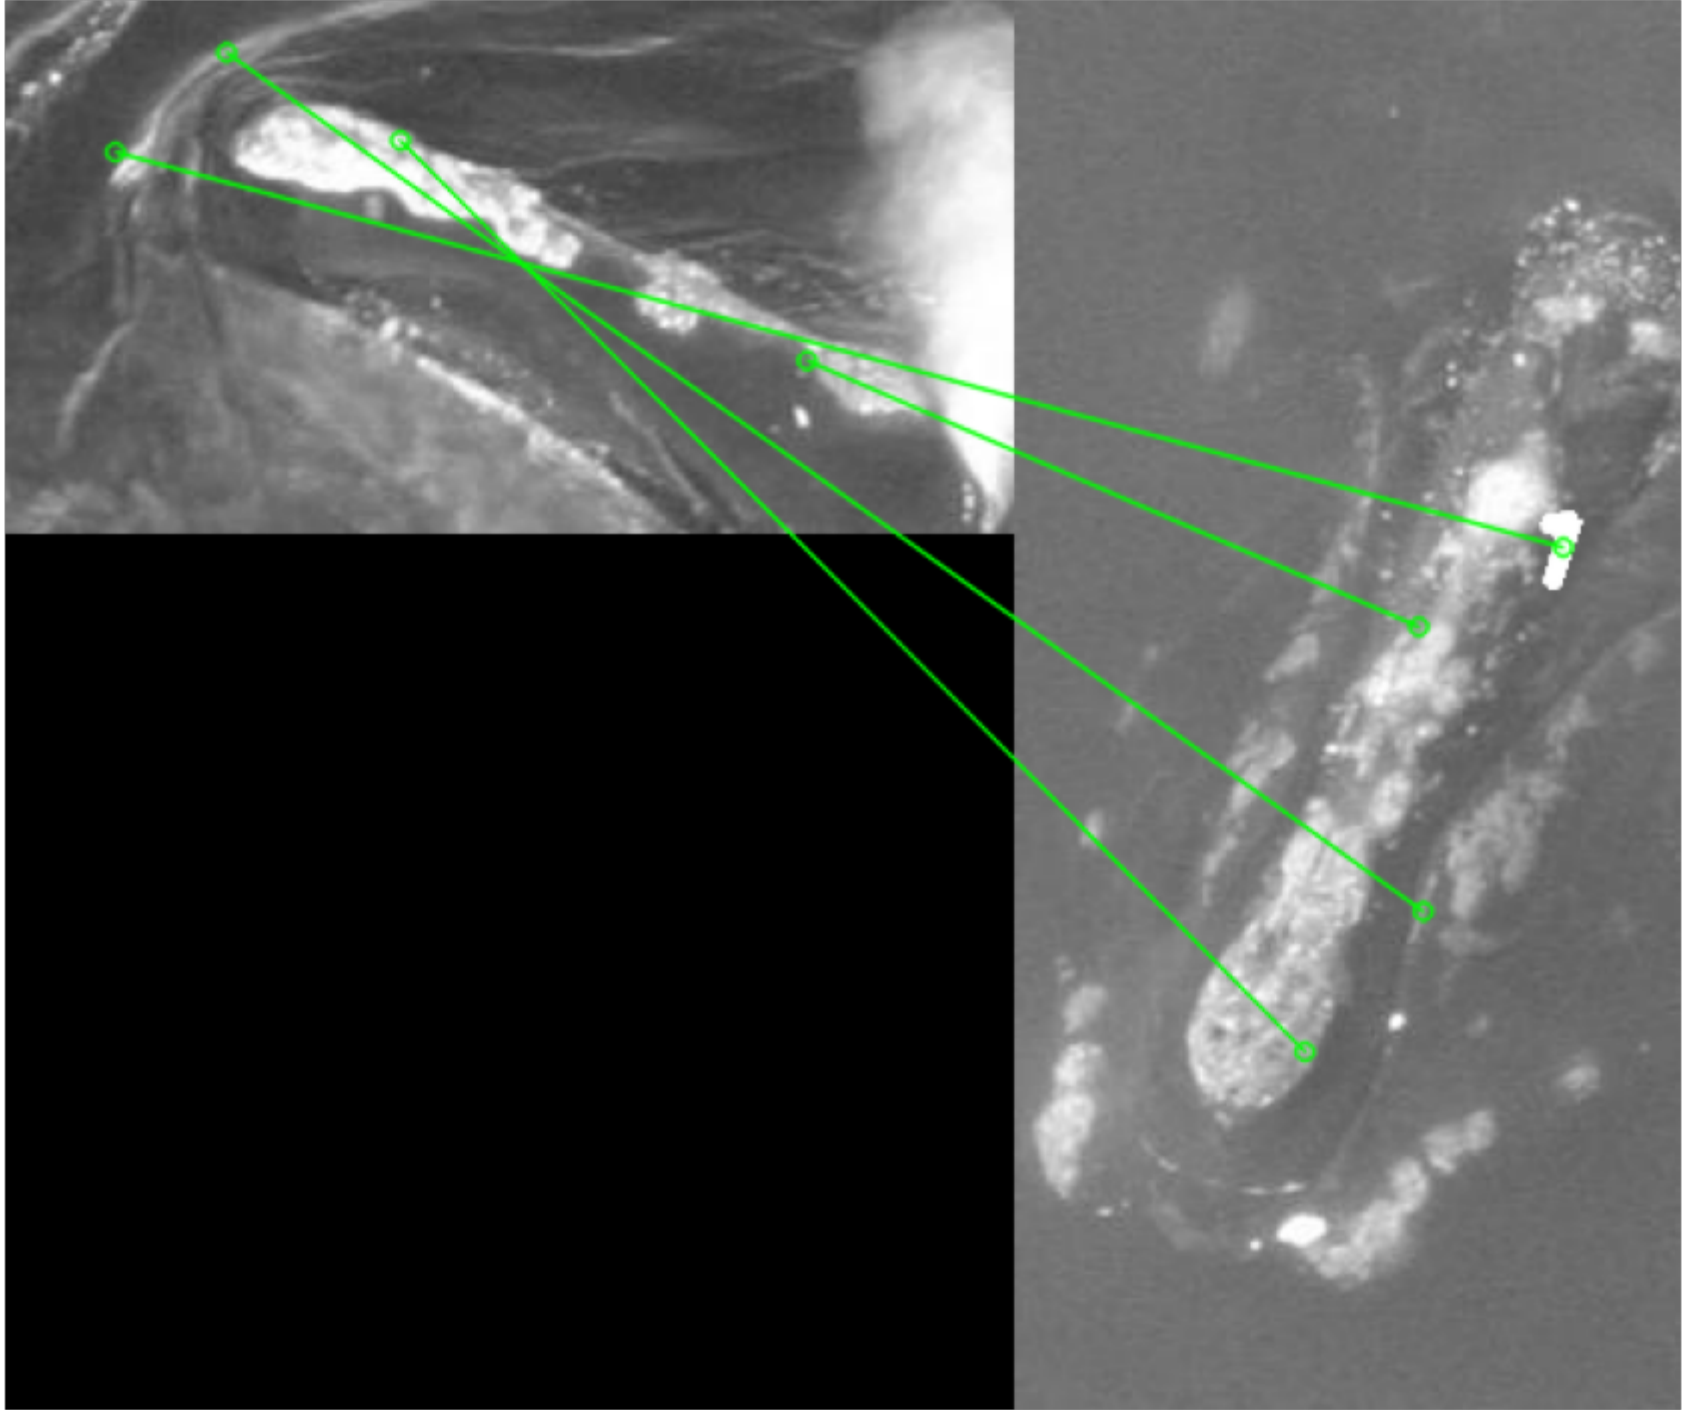
```
In [276]: 1 index_params = dict(algorithm=1, trees=5)
2 search_params = dict(checks=50)
3 flann = cv.FlannBasedMatcher(index_params, search_params)
4 matches = flann.knnMatch(des1, des2, k=2)
```

```
In [277]: 1 # store all the good matches as per Lowe's ratio test.
2 good = []
3 for m, n in matches:
4     if m.distance < 0.7*n.distance:
5         good.append(m)
6     if len(good) >= MIN_MATCH_COUNT:
7         src_pts = np.float32([kp1[m.queryIdx].pt for m in good]).reshape(-1, 1, 2)
8         dst_pts = np.float32([kp2[m.trainIdx].pt for m in good]).reshape(-1, 1, 2)
9         M, mask = cv.findHomography(src_pts, dst_pts, cv.RANSAC, 5.0)
10        matchesMask = mask.ravel().tolist()
11        h, w = img1.shape
12        pts = np.float32([[0, 0], [0, h - 1], [w - 1, h - 1], [w - 1, 0]]).reshape(-1, 1, 2)
13        dst = cv.perspectiveTransform(pts, M)
14        img2 = cv.polylines(img2, [np.int32(dst)], True, 255, 3, cv.LINE_AA)
15
```

```
In [278]: 1 draw_params = dict(matchColor=(0, 255, 0), # draw matches in green color
2                     singlePointColor=None,
3                     matchesMask=matchesMask, # draw only in-liers
4                     flags=2)
5 img3 = cv.drawMatches(img1, kp1, img2, kp2, good, None, **draw_params)
```



```
In [279]: 1 plt.figure(figsize=(15,15))
          2 plt.imshow(img3, 'gray')
          3 plt.axis('off')
          4 plt.show()
```



I found that sift features do not match well between two images of the same whale.

Experiment with saturation values

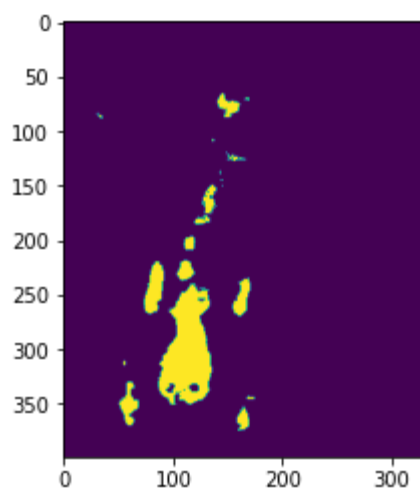
By experimening with Hue-Saturation Value (HSV), value of the images, I found a combination of values that helps detect the shape of the pattern on the head of the whale in a lot of images.

```
In [156]: 1 cropped_image = cv2.imread('/Users/pratikaheer/FALL22/CV/Whale/whale_samples/cropped/whale_48813/w_934.
```

```
In [151]: 1 hsv = cv2.cvtColor(cropped_image, cv2.COLOR_BGR2HSV)
```

```
In [154]: 1 h = hsv[:, :, 0]
          2 s = hsv[:, :, 1]
          3 v = hsv[:, :, 2]
          4
          5 HSV_MIN = np.array([1, 1, 1], np.uint8)
          6 HSV_MAX = np.array([50, 255, 255], np.uint8)
          7
          8 frame_threshold = cv2.inRange(hsv, HSV_MIN, HSV_MAX)
```

```
In [155]: 1 plt.imshow(frame_threshold)
          2 plt.show()
```



I could see a that HSV images were quite helpful to detect the shape of callousity patterns in a lot images. So I decide to use the HSV images for classifier to classify images to classes. So we convert all the cropped images to HSV images.

```
In [2]: 1 cropped_path = '/Users/pratikaheer/FALL22/CV/Whale/whale_samples/cropped/'
```

```
In [14]: 1 # convert cropped images to HSV images
          2 for more_dirs in os.listdir(cropped_path):
          3
          4     if more_dirs == '.DS_Store':
          5         continue
          6
          7     for file in os.listdir(cropped_path+more_dirs):
          8
          9         if file == '.DS_Store' or not file.endswith('.jpg'):
          10             continue
          11
          12         final_file = cropped_path+more_dirs+'/'+file
          13
          14
          15     # print(final_file)
          16     # read image
          17     read_cropped_image = cv2.imread(final_file)
          18
          19     # convert img to hsv
          20     hsv = cv2.cvtColor(read_cropped_image, cv2.COLOR_BGR2HSV)
          21     h = hsv[:, :, 0]
          22     s = hsv[:, :, 1]
          23     v = hsv[:, :, 2]
          24
          25     HSV_MIN = np.array([1, 5, 1], np.uint8)
          26     HSV_MAX = np.array([50, 255, 255], np.uint8)
          27
          28     # hsv_img = cv2.cvtColor(img, cv2.COLOR_BGR2HSV)
          29
          30     frame_threshed = cv2.inRange(hsv, HSV_MIN, HSV_MAX)
          31
          32     newpath = cropped_path+more_dirs+'/saturation/'
          33
          34     if not os.path.exists(newpath):
          35         os.makedirs(newpath)
          36
          37     cv2.imwrite(newpath+file, frame_threshed)
          38
          39
          40     # f, axarr = plt.subplots(1, 2, figsize=(10, 10))
          41     # plt.figure(figsize=(8, 8))
          42     # plt.imshow(read_cropped_image)
          43     # plt.imshow(frame_threshed)
          44
          45     # axarr[0].imshow(read_cropped_image)
          46     # axarr[1].imshow(frame_threshed)
          47
          48     # plt.imshow(h)
          49     # plt.imshow(cropped_image)
          50
          51     plt.axis('off')
          52     plt.show()
```

PCA Decomposition

```

In [7]: 1
2 PATH = '/Users/pratikaheer/FALL22/CV/Whale/whale_samples/cropped/'
3 N = 0
4 last_N = 0
5 xdim = 256
6 ydim = 256
7 D = np.zeros((140,xdim*ydim))
8 y_val = 0
9 y = []
10
11 whale_HM = {}
12
13 for foldername in os.listdir(PATH):
14     for file in os.listdir(PATH+foldername):
15         if( file.endswith(".jpg") ):
16
17             img = cv2.imread(PATH + foldername + "/" + file)
18             img = cv2.resize(img, (256, 256))
19
20             img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
21
22             sz = img.shape
23             # print(sz)
24             D[N,:] = img.flatten() # flatten image and pack into data matrix rows
25
26             if(False):
27                 plt.figure(figsize=(xdim/72,ydim/72))
28                 plt.imshow(img, cmap='gray') # display the image
29                 plt.axis('off')
30                 plt.show()
31                 sleep(0.01)
32                 clear_output(wait=True)
33             N = N + 1
34
35             y.append(y_val)
36
37         whale_HM[range(last_N, N+1)] = foldername
38         last_N = N
39         y_val += 1
40
41     print(foldername, N)
42
43 D = D.T # make images (num_images x pixels) in size
44 print( "loaded " + str(N) + " images" )

```

```

whale_48813 18
whale_65586 41
whale_36851 67
whale_38681 108
whale_28892 140
loaded 140 images

```

```

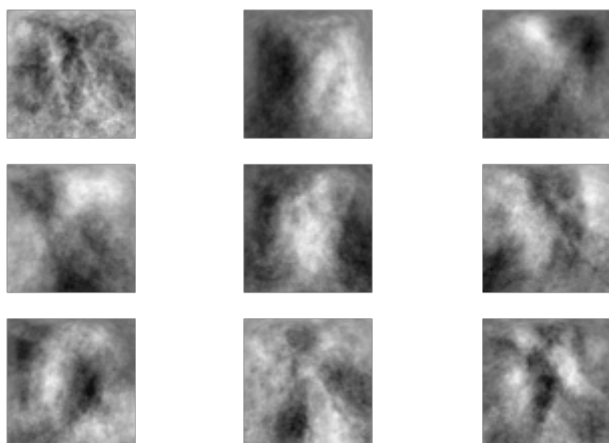
In [9]: 1 mu = np.mean(D, axis = 1)
2 mu = np.matlib.repmat(mu, N, 1)
3 D = D - mu.T
4 C2 = D.T@D;
5 val,vec = linalg.eig(C2);

```

```

In [10]: 1 for k in range(9):
2         e = D@vec[:,k]
3         plt.subplot(3,3,k+1);
4         plt.imshow( e.reshape(xdim,ydim), cmap='gray');
5         plt.axis('off')

```



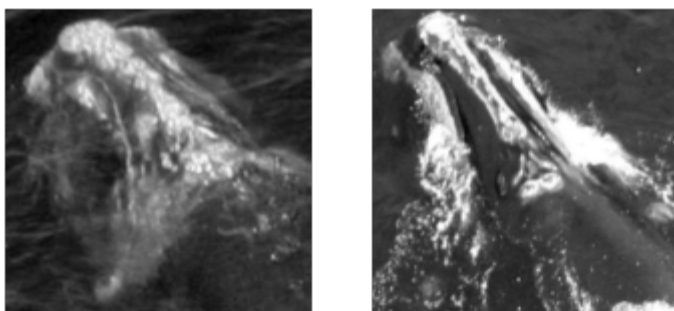

```
In [11]: 1 d = 3
2 P = np.zeros((d,N))
3 for k in range(d):
4     e = D@vec[:,k]
5     P[k,:] = e.T@D
```

```
In [23]: 1 def get_whale_from_photo_number(photo_num):
2
3     return [whale_HM[key] for key in whale_HM if photo_num in key][0]
```

```
In [13]: 1 Fnew = D + mu.T
2
3 j = 10
4 mindist = 1e10
5 for k in range(N):
6     if( j != k ):
7         dist = linalg.norm(P[:,k]-P[:,j])
8         if( dist < mindist ):
9             mindist = dist
10            minind = k
11
12 print(j, minind)
13 print(get_whale_from_photo_number(j), get_whale_from_photo_number(minind))
14
15 plt.subplot(121);
16 plt.imshow( Fnew[:,j].reshape(xdim,ydim), cmap='gray'); plt.axis('off');
17 plt.subplot(122);
18 plt.imshow( Fnew[:,minind].reshape(xdim,ydim), cmap='gray'); plt.axis('off')
```

```
10 27
whale_48813 whale_65586
```

```
Out[13]: (-0.5, 255.5, 255.5, -0.5)
```



PCA to classify images into categories

We use PCA projected images to find the closest image to a given image.

```
In [14]: 1 from collections import defaultdict
```

```
In [15]: 1 Fnew = D + mu.T
2
3 success = 0
4 HM = defaultdict(int)
5
6 for j in range(N):
7
8     mindist = 1e10
9     for k in range(N):
10         if( j != k ):
11             dist = linalg.norm(P[:,k]-P[:,j])
12             if( dist < mindist ):
13                 mindist = dist
14                 minind = k
15
16     if get_whale_from_photo_number(j) == get_whale_from_photo_number(minind):
17         success += 1
18
19 print("Accuracy", success/N)
```

```
Accuracy 0.2642857142857143
```

PCA does not give very good accuracy on the saturation images as the images are not will aligned and the whale faces are at a different position in the frame.

TSNE

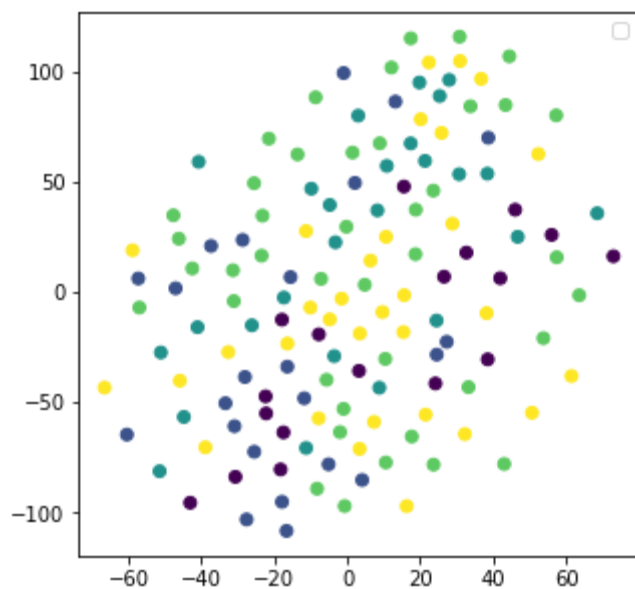
```
In [80]: 1 from matplotlib import cm
2 # import numpy as np
3 from sklearn import datasets
4 # from sklearn.manifold import TSNE
```

```
In [12]: 1 # TSNE (dimensionality reduction)
2 from sklearn.manifold import TSNE
3
4 tsne = TSNE(n_components=2, random_state=0)
5 X = tsne.fit_transform(D.T) # 2D representation
6 # y = range(len(digits.target_names)) # labels for visual
```

/Users/pratikaheer/opt/anaconda3/envs/cv/lib/python3.10/site-packages/sklearn/manifold/_t_sne.py:800: FutureWarning: The default initialization in TSNE will change from 'random' to 'pca' in 1.2.
warnings.warn(
/Users/pratikaheer/opt/anaconda3/envs/cv/lib/python3.10/site-packages/sklearn/manifold/_t_sne.py:810: FutureWarning: The default learning rate in TSNE will change from 200.0 to 'auto' in 1.2.
warnings.warn(

```
In [13]: 1 plt.figure(figsize=(5, 5))
2
3 plt.scatter(X[:, 0], X[:, 1], c=y)
4
5 plt.legend()
6 plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



We can see some association in the graph. But overall, we can see that scatter plot is all over the place for TSNE. This is because, as with PCA, the faces of the whales are not perfectly aligned.

Classification

SVM classifier

We use SVM on the pixel value of HSV images to classify images into respective categories.

```
In [1]: 1 from skimage.transform import resize
2 from sklearn.metrics import accuracy_score
3 from skimage.io import imread
```

```
In [2]: 1 cropped_path = '/Users/pratikaheer/FALL22/CV/Whale/whale_samples/cropped/'
```

```
In [5]: 1 flat_data_arr = []
2 target_arr = []
3
4 category_index = 0
5
6 for more_dirs in os.listdir(cropped_path):
7
8     for file in os.listdir(cropped_path+more_dirs+"/saturation"):
9
10         newpath = cropped_path+more_dirs + "/" + file
11         img_array = imread(newpath)
12         img_resized=resize(img_array,(256,256,3))
13         flat_data_arr.append(img_resized.flatten())
14         target_arr.append(category_index)
15
16     category_index += 1
```

```
In [6]: 1 flat_data=np.array(flat_data_arr)
        2 target=np.array(target_arr)
        3 df=pd.DataFrame(flat_data)
        4 df['Target']=target
        5 x=df.iloc[:, :-1]
        6 y=df.iloc[:, -1]
```

```
In [7]: 1 from sklearn import svm
        2 from sklearn.model_selection import GridSearchCV
        3 param_grid={'C':[0.1,1,10,100], 'gamma':[0.0001,0.001,0.1,1], 'kernel':['rbf', 'poly']}
        4 svc = svm.SVC(probability=True)
        5 model = GridSearchCV(svc,param_grid)
```

```
In [8]: 1 from sklearn.model_selection import train_test_split
        2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20,random_state=77,stratify=y, shuffle=
```

I am using grid search to identify the best parameters for the SVM model.

```
In [9]: 1 model.fit(x_train,y_train)
```

```
Out[9]:  ▸ GridSearchCV
        ▸ estimator: SVC
           ▸ SVC
```

```
In [10]: 1 # Accuracy on the train set
        2 accuracy_score(model.predict(x_train), y_train)
```

```
Out[10]: 1.0
```

We see that the model is overfitting because the accuracy is 100% on the training set.

```
In [11]: 1 y_pred=model.predict(x_test)
```

```
In [12]: 1 print(f"The model is {accuracy_score(y_pred,y_test)*100}% accurate")
```

The model is 28.57142857142857% accurate

Let us try to see how the accuracy is with just two classes. We take images of two whales and name the whales alex and julia. We create a classifier that trains on the HSV values of two whales and use SVM classifier for testing.

```
In [25]: 1 import pandas as pd
        2 import os
        3 from skimage.transform import resize
        4 from sklearn.metrics import accuracy_score
        5 from skimage.io import imread
        6 import numpy as np
        7 import matplotlib.pyplot as plt
        8
        9 Categories = ['alex', 'julia']
       10
       11 datadir = '/Users/pratika/her/FALL22/CV/Whale/investigation/24458/outputs'
       12
       13 flat_data_arr = []
       14 target_arr = []
       15
       16 for category in Categories:
       17     print(f'loading... category : {category}')
       18     path=os.path.join(datadir,category)
       19     for img in os.listdir(path):
       20         img_array = imread(os.path.join(path,img))
       21         img_resized=resize(img_array,(256,256,3))
       22         flat_data_arr.append(img_resized.flatten())
       23         target_arr.append(Categories.index(category))
```

```
loading... category : alex
loading... category : julia
```

```
In [26]: 1 flat_data=np.array(flat_data_arr)
        2 target=np.array(target_arr)
        3 df=pd.DataFrame(flat_data)
        4 df['Target']=target
        5 x=df.iloc[:, :-1]
        6 y=df.iloc[:, -1]
```



```
In [27]: 1 from sklearn import svm
2 from sklearn.model_selection import GridSearchCV
3 param_grid={'C':[0.1,1,10,100], 'gamma':[0.0001,0.001,0.1,1], 'kernel':['rbf', 'poly']}
4 svc = svm.SVC(probability=True)
5 model = GridSearchCV(svc,param_grid)
```

```
In [28]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20,random_state=77,stratify=y)
```

```
In [29]: 1 model.fit(x_train,y_train)
```

```
Out[29]: ▸ GridSearchCV
          ▸ estimator: SVC
              ▸ SVC
```

```
In [7]: 1 y_pred=model.predict(x_test)
```

```
In [8]: 1 print("The predicted Data is :")
2 print(y_pred)
3 print("The actual data is:")
4 print(np.array(y_test))
5 print(f"The model is {accuracy_score(y_pred,y_test)*100}% accurate")
```

The predicted Data is :
 [1 1 1 1 1 1]
 The actual data is:
 [1 1 0 1 1 0]
 The model is 71.42857142857143% accurate

Image Transformers

CLIP stands for Contrastive Language-Image Pre-Training is trained on ImageNet dataset and can be used to build a classifier by training it with the set of our images. It is preferred for zero shot learning, as it can classify the things it has never seen before.

```
In [ ]: 1 from sentence_transformers import SentenceTransformer, util
2 from PIL import Image
3 import glob
4 import os
5
6 # Load the OpenAI CLIP Model
7 print('Loading CLIP Model...')
8 model = SentenceTransformer('clip-ViT-B-32')
```

/Users/pratikaheer/opt/anaconda3/envs/cv/lib/python3.10/site-packages/tqdm/auto.py:22: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html (https://ipywidgets.readthedocs.io/en/stable/user_install.html)
 from .autonotebook import tqdm as notebook_tqdm

Loading CLIP Model...

```
In [ ]: 1 cropped_path = '/Users/pratikaheer/FALL22/CV/Whale/whale_samples/cropped/'
2
3 image_names = []
4
5 for more_dirs in os.listdir(cropped_path):
6
7     if more_dirs == '.DS_Store':
8         continue
9
10    image_names.extend(list(glob.glob(cropped_path+ more_dirs +'/*.jpg')))
```

```
In [ ]: 1 image_data = list(zip(image_names, y))
```

```
In [47]: 1 import random
2 random.shuffle(image_data)
3
4 test_size = int(len(image_data) * 0.2)
5 valid_size = int(len(image_data) * 0.1)
6 train_size = len(image_data) - test_size - valid_size
7
8
9 train_data = image_data[:train_size]
10 valid_data = image_data[train_size:train_size+valid_size]
11 test_data = image_data[-test_size:]
```

```
In [48]: 1 print("Images:", len(train_data))
2 encoded_images_train = model.encode([Image.open(filepath[0]) for filepath in train_data], batch_size=128)
```

Images: 98

Batches: 100% | 1/1
[00:06<00:00, 6.30s/it]

```
In [49]: 1 encoded_images_test = model.encode([Image.open(filepath[0]) for filepath in test_data], batch_size=128)
```

Batches: 100% | 1/1
[00:01<00:00, 1.97s/it]

Once we have the embeddings from CLIP model, we can use K-nearest neighbors classifier to classify the image embeddings obtained from the transformer.

```
In [50]: 1 from sklearn.neighbors import KNeighborsClassifier
2
3 neigh = KNeighborsClassifier(n_neighbors=3)
4 neigh.fit(encoded_images_train, [e[1] for e in train_data])
```

```
Out[50]: KNeighborsClassifier
KNeighborsClassifier(n_neighbors=3)
```

I experimented with various k values but 3 seems to give best results.

```
In [51]: 1 success = 0
2 for i in range(len(test_data)):
3
4     test_prediction = neigh.predict(encoded_images_test[i].unsqueeze(0))
5
6     if test_prediction[0] == test_data[i][1]:
7
8         success += 1
9
10 print("Accuracy", success/len(test_data))
```

Accuracy 0.5357142857142857

Limitations

Neural network based algorithms like CNN should work well to detect the callosity patterns of the whales. We can also experiment with different filters. It is really difficult to separate the features of the head of the whale from the water and the glare of the sun. HSV images are the closest I could come to getting the head patterns. With better object localization, it is possible that we can get the features of the head more clearly. I could also have expanded the labeled image set using some augmentation techniques. For every image, we would run it through some filters which are meant to signify the variance between the images of the set, and then use filter outputs as additional training data. Some of the filters we could use low pass filter (smoothing), high pass filter (edge detection or sharpening), and various affine transform such as rotation, scaling, and pixel-wise shift. It would be also interesting to see if we can use ensemble models each focusing on different features from pre-processing (e.g edge detection and HOG values).

Conclusion

I can see good accuracy with just two categories. When I include more categories their accuracy goes down as the difference between categories is very little because there are only ~30 examples to train from. This is in line with the existing methods, where others have also could achieve accuracy in the range of 50-65 percent. With classification methods like SVM, I can see that the model is overfitting to the training data. I conclude that, while I got reasonable results for two classes. It would be really useful to train a model on whale images with sufficient pre-processing techniques.